

Homework #6

Algorithms I

600.463

Spring 2017

Due on: Thursday, April 6th, 11:59pm

Late submissions: will NOT be accepted

Format: Please start each problem on a new page.

Where to submit: On Gradescope under HW6

Please type your answers; handwritten assignments will not be accepted.

To get full credit, your answers must be explained clearly,
with enough details and rigorous proofs.

April 17, 2017

Problem 1 (20 points)

Suppose you are given an undirected graph G with weighted edges and a minimum spanning tree T of G .

- Design an algorithm to update the minimum spanning tree when the weight of a single edge e is *increased*.
- Design an algorithm to update the minimum spanning tree when the weight of a single edge e is *decreased*.

In both cases, the input to your algorithm is the edge e and its new weight; your algorithms should modify T so that it is still a minimum spanning tree. Analyze the running time of your algorithms and prove the correctness.

Answer:

- Consider the case that edge e is *increased*: if e is not an edge in T , no need to change anything, since e cannot be a part of the MST by increasing its weight. If add e to T , the new T creates a subgraph H with one cycle. Then

e is the longest edge in that cycle even before we increase its weight. Thus increasing the weight of e only makes it more useless.

If e is an edge of T , we can first delete e from T , which creates an intermediate spanning forest F with two components. Let e_{min} be the minimum-weight edge with one endpoint in each component of F . So the new MST is $F \cup e_{min}$. (It is possible that e and e_{min} are the same edge.)

Then we need a method to find this e_{min} . We first mark all the vertices in one component of F , by using BFS/DFS in $O(V + E)$ time. Then we go through every edge of G and find the minimum-weight edge with exactly one marked endpoint in $O(E)$ time.

- Consider the case that edge e is *decreased*: if e is an edge in T , no need to change anything, since e still belongs to the MST after we decrease its weight. If deleting e from T , there obtains an intermediate spanning forest F . Then e is a safe edge with respect to F even before we decrease its weight. Thus decreasing the weight of e only makes it safer.

If e is not an edge of T , the graph $U = T \cup \{e\}$ must contain a cycle, since there is a unique path in T between the endpoints of e . Let e_{max} be the maximum-weight edge in this cycle. The new MST is $T' = U \setminus \{e_{max}\}$. (e and e_{max} could be the same edge.)

We can use BFS/DFS to find e_{max} in U in $O(V + E)$ time.

Problem 2 (20 points)

Problem 2.1 (10 points)

Let $G = (V, E)$ be a directed, weighted graph with weight function $w : E \rightarrow \mathbb{R}$. Give an $O(VE)$ -time algorithm to find, for each vertex $v \in V$, the value $\delta^*(v) = \min_{u \in V} \{\delta(u, v)\}$, where $\delta(u, v)$ is the *shortest-path weight* from u to v defined in the textbook.

Answer:

You can do this by making a new dynamic-programming algorithm that probably looks a bit like Bellman-Ford, with the updates going in the “opposite” direction. Let’s consider a solution without a new algorithm at all. Construct a graph $G' = (V', E')$, where $V' = V \cup \{s\}$ and $E' = E \cup \{(s, u) | u \in V\}$, i.e. we have a new vertex s and there is an edge from the new vertex to every vertex in the graph. Set the weight of these new edges to be $w(s, u) = 0$. Then, just run Bellman-Ford algorithm on G' starting from s .

Claim 0.1. $\delta_G^*(v) = \delta_{G'}(s, v)$

Proof. If we want to prove this, we should show that $\delta_G^*(v) \leq \delta_{G'}(s, v)$ and $\delta_G^*(v) \geq \delta_{G'}(s, v)$. To show the former, suppose $s \rightarrow u \rightsquigarrow v$ is a shortest path from s to v in G' , then $w(s \rightarrow u \rightsquigarrow v) = w(u \rightsquigarrow v)$, and $u \rightsquigarrow v$ is a shortest path from u to v in G . Thus, we have $\delta_G^*(v) \leq \delta_G(u, v) = w(s \rightarrow u \rightsquigarrow v) = \delta_{G'}(s, v)$. To argue the other direction, suppose that $u = \operatorname{argmin}_{u \in V} \{\delta_G(u, v)\}$, i.e., that $\delta_G(u, v) = \delta_G^*(v)$. Then a valid path from s to v in G' is $s \rightarrow u \rightsquigarrow v$ with weight $\delta_G^*(v)$, and we conclude that $\delta_{G'}(s, v) \leq \delta_G(u, v) = \delta_G^*(v)$. \square

Problem 2.2 (10 points)

Let $G = (V, E)$ be a directed, weighted graph with nonnegative weight function $w : E \rightarrow \{1, 2, \dots, W\}$ for some nonnegative integer W . Devise an algorithm to compute the shortest path distances from a given source vertex s in $O(WV + E)$ time.

Answer:

Let's consider two types of solutions:

- **Solution 1:** Run Dijkstra's algorithm but with a priority queue that is geared specifically towards this problem. In particular, at any time, only W different distances are in the queue, with at most WV distance in total. You could thus implement the priority queue as follows: create an array of size WV containing all possible distances. Keep a pointer to the current min distance. Each array entry contains (a pointer to) an unordered list of elements having that distance. To *EXTRACT-MIN*, remove any element from the current list in $O(1)$ time. If that list is empty, advance the pointer until finding a nonempty list. There are $O(WV)$ pointer advances in total (amortized analysis here). To *DECREASE-KEY*, simply move the element from one list to another in $O(1)$ time, splicing it out of one and inserting it at the front of the other. So the total cost is $O(WV + V + E)$ for V *EXTRACT-MIN*s and E *DECREASE-KEY*s. Note that this is *not* a general-purpose priority queue — the decrease key is only allowed to decrease as low as the current value of the pointer, which is sufficient for Dijkstra's algorithm.
- **Solution 2:** Reduce a graph with weights in $\{1, 2, \dots, W\}$ to a graph with edges of weights all 1 with the same shortest path distances. BFS finds shortest paths in unit-weighted graphs. So you can apply BFS here at that point. Specifically, the new graph is as follows. For each u , create W copies u_1, u_2, \dots, u_W . That is

$$V' = \{u_i | u \in V, 1 \leq i \leq W\}. \quad (1)$$

As for the edges, string together all u_i in a chain, i.e., $u_1 \rightarrow u_2 \rightarrow \dots \rightarrow u_W$. Now the idea is that we can mimic the original edges by being careful about where we put them. All edges pointed towards v originally should target v_1 in the new graph. An edge (u, v) pointed out of u should originate from the $w(u, v)$ th copy of u (i.e., $u_{w(u,v)}$). Thus, following the edge (u, v) in the original graph corresponds to the $w(u, v)$ -hop path $u_1 \rightarrow u_2 \rightarrow \dots \rightarrow u_{w(u,v)} \rightarrow v_1$ in the augmented graph. We have

$$E' = \{(u_w, v_1) \mid (u, v) \in E, w = w(u, v)\} \cup \{(u_i, u_{i+1})\} \quad (2)$$

The new graph has $|V'| = W|V|$ vertices and $|E'| = |E| + (W - 1)|V|$ edges. Applying BFS thus runs in $O(WV + E)$ time.

Thinking along the lines of the second solution, you may be tempted to create $w(u, v)$ vertices for each edge (u, v) . The problem with that is that you would have $\Theta(W^2E)$ edges in the worst case, which is worse than $O(WV + E)$. So it is important that we should be careful in the construction to avoid this. Thus, we use a single chain of W vertices per original vertex.