

Homework #5 Solutions  
Introduction to Algorithms  
601.433/633  
Spring 2020

**Due on:** Tuesday, March 24th, 12pm

**Where to submit:** On Gradescope, please mark the pages for each question

## 1 Problem 1 (25 points)

### 1.1 Problem 1.1 (10 points)

Suppose we wish not only to increment a counter of length  $m \in \mathbb{N}$  but also to reset it to zero (i.e., make all bits in it 0). Counting the time to examine or modify a bit as  $\Theta(1)$ , show how to implement INCREMENT and RESET operations on a counter (represented as an array of bits) so that *any* sequence of  $n$  operations takes  $O(1)$  amortized time per operation.

You may assume that the counter is initially zero and that you may access and modify any specific bit (say bit  $j \in [m]$ ) in  $O(1)$  time. Additionally, you may assume that the total count in the counter never exceeds  $2^m - 1$  during the course of the  $n$  operations.

Prove correctness and running time of your algorithms.

(Hint: Keep a pointer to the highest-order 1.)

*Proof.* We introduce a new field  $max[A]$  to hold the index of the high-order 1 in  $A$ . Initially,  $max[A]$  is set to 1, since the low-order bit of  $A$  is at index 0, and there are initially no 1's in  $A$ . The value of  $max[A]$  is updated as appropriate when the counter is incremented or reset, and we use this value to limit how much of  $A$  must

be looked at to reset it. By controlling the cost of *RESET* in this way, we can limit it to an amount that can be covered by credit from earlier *INCREMENT*s.

---

```

INCREMENT(A)
i ← 0
while i < length[A] and A[i] = 1 do
    A[i] ← 0
    i ← i + 1
    if i < length[A] then
        A[i] ← 1
        if i > max[A] then
            max[A] ← i
        end if
    else
        max[A] ← -1
    end if
end while

```

---



---

```

RESET(A)
for i ← 0 to max[A] do
    A[i] ← 0
end for
max[A] ← -1

```

---

As for the counter in the book, we assume that it costs \$1 to flip a bit. In addition, we assume it costs \$1 to update  $\text{max}[A]$ .

Setting and resetting of bits by *INCREMENT* will work exactly as for the original counter in the book: \$1 will pay to set one bit to 1; \$1 will be placed on the bit that is set to 1 as credit; the credit on each 1 bit will pay to reset the bit during incrementing.

In addition, we will use \$1 to pay to update  $\text{max}$ , and if  $\text{max}$  increases, we will place an additional \$1 of credit on the new high-order 1. (If  $\text{max}$  doesn't increase, we can just waste that \$1— it won't be needed.) Since *RESET* manipulates bits at positions only up to  $\text{max}[A]$ , and since each bit up to there must have become the high-order 1 at some time before the high-order 1 got up to  $\text{max}[A]$ , every bit seen by *RESET* has \$1 credit on it. So the zeroing of bits of  $A$  by *RESET* can be completely paid for by the credit stored on the bits. We just need \$1 to pay for resetting  $\text{max}$ .

Thus, charging \$4 for each *INCREMENT* and \$1 for each *RESET* is sufficient, so the sequence of  $n$  *INCREMENT* and *RESET* operations takes  $O(n)$  time.  $\square$

## 1.2 Problem 1.2 (15 points)

Design a data structure to support the following two operations for a set  $S$  of integers, which allows duplicate values:

- $\text{INSERT}(S, x)$  inserts  $x$  into  $S$ .
- $\text{DELETE-LARGER-HALF}(S)$  deletes the largest  $\lceil |S|/2 \rceil$  elements from  $S$ .

Explain how to implement this data structure so that any sequence of  $m$   $\text{INSERT}$  and  $\text{DELETE-LARGER-HALF}$  operations runs in amortized  $O(1)$  time per operation. Your implementation should also include a way to output the elements of  $S$  in  $O(|S|)$  time.

Prove the running time of your implementation.

*Proof.* We can use a simple array  $A$  to be our underlying data structure. Array  $A$  contains a set  $S$  of integers. The operations are as follows.

- $\text{INSERT}(S, x)$ : The element can be inserted to the last element in  $A$ .
- $\text{DELETE-LARGER-HALF}(S)$ : First, find the median  $mid$  of  $S$  using array  $A$ . We then partition  $S$  into two subsets,  $S_1$  and  $S_2$ , where all the elements in  $S_1$  are less than  $mid$  and the elements in  $S_2$  are larger than and equal to  $mid$  by comparing each element in  $A$  with  $mid$ . Last, we delete  $S_2$  and set  $S = S_1$ .

Suppose there are  $m$  operations performed on  $S$ . For each  $i = 1, 2, \dots, m$ , let  $c_i$  be the actual cost of the  $i$ th operation and  $A_i$  be the data structure that results after applying the  $i$ th operation to data structure  $A_{i-1}$ . Since we use the linear time algorithm for finding the median, the running time of the algorithm for discarding the larger half is  $O(m)$  for an  $m$ -element array. Suppose the actual time for discarding

the larger half is  $km$  for some constant  $k$ . We define the *potential function*  $\Phi$  on each data structure  $A_i$  to be

$$\Phi(A_i) = 2k \times |A_i|,$$

where  $|A_i|$  is the number of elements in array  $A_i$ . For an operation  $i$ , we analyze the cost by considering the following two cases.

*INSERT*( $S, x$ ): The actual cost  $c_i$  for insertion is 1. Then, the amortized cost

$$\hat{c}_i = c_i + \Phi(A_i) - \Phi(A_{i-1}) = 1 + 2k \times |A_i| - 2k \times |A_{i-1}| = 1 + 2k,$$

since  $|A_i| = |A_{i-1}| + 1$  when inserting an element to  $A_{i-1}$ . The amortized cost therefore is  $O(1)$ .

*DELETE-LARGER-HALF*( $S$ ): As mentioned above, the actual cost for deleting the larger half is  $km$  if there are  $m$  elements. Then, the amortized cost

$$\begin{aligned} \hat{c}_i &= c_i + \Phi(A_i) - \Phi(A_{i-1}) \\ &= k \times |A_{i-1}| + 2k \times |A_i| - 2k \times |A_{i-1}| \\ &= k \times |A_{i-1}| + 2k \times \frac{1}{2}|A_{i-1}| - 2k \times |A_{i-1}| \\ &= k \times |A_{i-1}| + k \times |A_{i-1}| - 2k \times |A_{i-1}| = 0 \end{aligned}$$

since  $|A_i| = \frac{1}{2}|A_{i-1}|$  when deleting the larger half of  $A_{i-1}$ . The amortized cost therefore is 0.

Therefore, the  $m$  operations run in  $O(m)$  time. □

## Problem 2 (10 points)

Let  $G = (V, E)$  be a directed graph.  $a \in V$  is a *central* vertex if for all  $b \in V$  there exists a path from  $a$  to  $b$ . Provide an  $O(|V| + |E|)$  time algorithm to test whether graph  $G$  has a central vertex.

Prove the correctness of your algorithm and analyze the running time.

*Proof.* First, we notice that the directed graph  $G$  can be cyclic. Is there any method we can get rid of all the cycles? Let's consider one possible algorithm described as following.

1. Use Kosaraju's algorithm to find all strongly connected components of the given graph  $G$ .
2. Build a new graph  $G^*$  by shrinking each SCC of  $G$  into one representative vertex and keeping the edges between SCCs.
3. If there is more than one vertex that has no incoming edges, return FALSE; else, return TRUE.

**Running Time:** Assuming that the graph is presented by adjacency list, Kosaraju's algorithm runs in  $O(|V| + |E|)$  time. Building the new graph  $G^*$  and finding outgoing-only vertices take  $O(|V| + |E|)$ . In total, the time complexity is  $O(|V| + |E|)$ .

**Correctness:** I eliminated some details here. Briefly, let's consider the graph  $G^*$ . Clearly,  $G^*$  is a DAG. DAG has at least one vertex with no incoming edges (we can use proof by contradiction, eliminated here). So the only case left, one vertex with no incoming edges is actually the central vertex. This is because all other vertices have at least one incoming edge and can be reached by one of vertices. If there are more than one such vertex, e.g. two vertices, one cannot be reached by the other, which contracts the definition of central vertex.  $\square$

## 2 Problem 3 (15 points)

You're helping some analysts monitor a collection of networked computers, tracking the spread of fake information. There are  $n$  computers in the system, labeled  $C_1, C_2, \dots, C_n$ , and as input you're given a collection of trace data indicating the times at which pairs of computers communicated. Thus the data is a sequence of ordered triples  $(C_i, C_j, t_k)$ ; such a triple indicates that  $C_i$  and  $C_j$  exchanged bits at time  $t_k$ . There are  $m$  triples total.

We'll assume that the triples are presented to you in sorted order of time. For purposes of simplicity, we'll assume that each pair of computers communicates at most once during the interval you're observing. The analysts you're working with

would like to be able to answer questions of the following form: If the fake information was generated by computer  $C_a$  at time  $x$ , could it possibly have been sent to  $C_b$  by time  $y$ ? The mechanics of communicating the information are simple: if a computer containing the fake information  $C_i$  communicates with another computer  $C_j$  that hasn't received that information yet by time  $t_k$  (in other words, if one of the triples  $(C_i, C_j, t_k)$  or  $(C_j, C_i, t_k)$  appears in the trace data), then computer  $C_j$  receives the fake information, starting at time  $t_k$ .

The fake information can thus spread from one machine to another across a sequence of communications, provided that no step in this sequence involves a move backward in time. Thus, for example, if  $C_i$  has received the fake information by time  $t_k$ , and the trace data contains triples  $(C_i, C_j, t_k)$  and  $(C_j, C_q, t_r)$ , where  $t_k \leq t_r$ , then  $C_q$  will receive the fake information via  $C_j$ . (Note that it is okay for  $t_k$  to be equal to  $t_r$ ; this would mean that  $C_j$  had open connections to both  $C_i$  and  $C_q$  at the same time, and so the information would have been sent from  $C_i$  to  $C_q$ .)

Design an algorithm that answers questions of this type: given a collection of trace data, the algorithm should decide whether the fake information generated by computer  $C_a$  at time  $x$  could have been received by computer  $C_b$  by time  $y$ . The algorithm should run in time  $O(m + n)$ .

Prove correctness and running time as usual.

*Proof.* We will first create a graph  $G = (V, E)$ . For each element  $(C_i, C_j, t_k)$  in the trace data we create vertices  $(C_i, t_k)$  and  $(C_j, t_k)$  and add directed edges in both directions between the vertices. Additionally, we add an edge from  $(C_i, t')$  to  $(C_i, t_k)$  where  $t'$  is the preceding time step (the previously last one)  $C_i$  appeared in. We do the same for  $C_j$ .

Constructing this graph takes  $O(m)$  time by maintaining an array of linked lists – a linked list of time steps (in descending order) associated with each machine  $C_i$ .

Given a pair of machines  $C_a, C_b$  and a pair of time steps  $x, y$  to check whether the fake information introduced at time  $x$  in machine  $C_a$  could have been received by machine  $C_b$  by time  $y$  – we run DFS from every vertex  $(C_a, t)$  such that  $t \geq x$  is the smallest time step larger than  $x$  that appears in the linked list for  $C_a$ . If we reach a vertex  $(C_b, t')$  such that  $t' \leq y$  then we return YES else we return NO.

**Correctness** If the algorithm returns YES, then there exists a path from  $C_a$  to  $C_b$  in increasing order of time starting at a time after  $x$  and ending at a time before  $y$ . Hence, by the rules by which fake information spreads,  $C_b$  could have received the information.

If the algorithm returns NO – assume for contradiction that  $C_a$  could have sent the information to  $C_b$  by time  $y$  if  $C_a$  received the information at time  $x$ . Then, by how we constructed the graph, there must be a sequence of machines with open connections in non-decreasing order of time. It is easy to show that there must be a directed path from a vertex  $(C_a, t)$  (for  $t \geq x$ ) to another  $(C_b, t')$  (for  $t' \leq y$ ) in  $G$ . This is a contradiction to the fact the algorithm returned NO.

**Running Time** The running time bound follows by showing that the time to create  $G$  is  $O(m)$  when the trace data is presented in sorted order of time and running DFS requires  $O(m + n)$  time.

□